

```

#Importing Libraries
# please do go through this python notebook:
import warnings
warnings.filterwarnings("ignore")

import csv
import pandas as pd#pandas to create small dataframes
import datetime #Convert to unix time
import time #Convert to unix time
# if numpy is not installed already : pip3 install numpy
import numpy as np#Do arithmetic operations on arrays
# matplotlib: used to plot graphs
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns#Plots
from matplotlib import rcParams#Size of plots
from sklearn.cluster import MiniBatchKMeans, KMeans#Clustering
import math
import pickle
import os
# to install xgboost: pip3 install xgboost
#!pip3 install xgboost
import xgboost as xgb

import warnings
import networkx as nx
import pdb
import pickle
from pandas import HDFStore, DataFrame
from pandas import read_hdf
from scipy.sparse.linalg import svds, eigs
import gc
from tqdm import tqdm
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score

#!pip3 install wget

#!wget --header="Host: doc-0o-bk-docs.googleusercontent.com" --
header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64)
AppleWebKit/537.36 (KHTML, like Gecko) Chrome/90.0.4430.212
Safari/537.36" --header="Accept:
text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image
/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --
header="Accept-Language: en-US,en;q=0.9" --header="Cookie:
AUTH_nso6dcnlmbidkt5qr539a2jiefc09pqv_nonce=iak2ig7rpq664" --
header="Connection: keep-alive" "https://doc-0o-bk-
docs.googleusercontent.com/docs/securesc/nss2f5s2soorprev6d4t4qp3n5ekp
9nh/

```

```
evl2j2j4t5hronicnhsbdlslbnbl9qk3/1622116650000/06629147635963609455/13
017565264516993811/1fDJptLCFEWNV5UNGPc4geTykgFI3PDCV?
e=download&authuser=0&nonce=iak2ig7rpq664&user=13017565264516993811&ha
sh=fvl5s6dohfnqle6k8q3koe9jr2mhe6jr" -c -0 'storage_sample_stage4.h5'
```

```
#reading
```

```
from pandas import read_hdf
df_final_train = read_hdf('storage_sample_stage4.h5',
    'train_df',mode='r')
df_final_test = read_hdf('storage_sample_stage4.h5',
    'test_df',mode='r')
```

```
#df_final_train.columns
```

```
y_train = df_final_train.indicator_link
y_test = df_final_test.indicator_link
```

```
df_final_train.drop(['source_node',
    'destination_node','indicator_link'],axis=1,inplace=True)
df_final_test.drop(['source_node',
    'destination_node','indicator_link'],axis=1,inplace=True)
```

```
estimators = [10,50,100,250,450]
```

```
train_scores = []
```

```
test_scores = []
```

```
for i in estimators:
```

```
    clf = RandomForestClassifier(bootstrap=True, class_weight=None,
    criterion='gini',
```

```
        max_depth=5, max_features='auto', max_leaf_nodes=None,
```

```
        min_impurity_decrease=0.0, min_impurity_split=None,
```

```
        min_samples_leaf=52, min_samples_split=120,
```

```
        min_weight_fraction_leaf=0.0, n_estimators=i, n_jobs=-
```

```
1,random_state=25,verbose=0,warm_start=False)
```

```
    clf.fit(df_final_train,y_train)
```

```
    train_sc = f1_score(y_train,clf.predict(df_final_train))
```

```
    test_sc = f1_score(y_test,clf.predict(df_final_test))
```

```
    test_scores.append(test_sc)
```

```
    train_scores.append(train_sc)
```

```
    print('Estimators = ',i,'Train Score',train_sc,'test
```

```
Score',test_sc)
```

```
plt.plot(estimators,train_scores,label='Train Score')
```

```
plt.plot(estimators,test_scores,label='Test Score')
```

```
plt.xlabel('Estimators')
```

```
plt.ylabel('Score')
```

```
plt.title('Estimators vs score at depth of 5')
```

```
Estimators = 10 Train Score 0.9063252121775113 test Score
0.8745605278006858
```

```
Estimators = 50 Train Score 0.9205725512208812 test Score
0.9125653355634538
```


```
Estimators = 100 Train Score 0.9238690848446947 test Score
```

0.9141199714153599

Estimators = 250 Train Score 0.9239789348046863 test Score
0.9188007232664732

Estimators = 450 Train Score 0.9237190618658074 test Score
0.9161507685828595


Text(0.5, 1.0, 'Estimators vs score at depth of 5')



```
depths = [3,9,11,15,20,35,50,70,130]
train_scores = []
test_scores = []
for i in depths:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None,
criterion='gini',
    max_depth=i, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=52, min_samples_split=120,
    min_weight_fraction_leaf=0.0, n_estimators=115, n_jobs=-
1,random_state=25,verbose=0,warm_start=False)
    clf.fit(df_final_train,y_train)
    train_sc = f1_score(y_train,clf.predict(df_final_train))
    test_sc = f1_score(y_test,clf.predict(df_final_test))
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('depth = ',i,'Train Score',train_sc,'test Score',test_sc)
plt.plot(depths,train_scores,label='Train Score')
plt.plot(depths,test_scores,label='Test Score')
plt.xlabel('Depth')
```

```
plt.ylabel('Score')
plt.title('Depth vs score at depth of 5 at estimators = 115')
plt.show()
```

```
depth = 3 Train Score 0.8916120853581238 test Score
0.8687934859875491
depth = 9 Train Score 0.9572226298198419 test Score
0.9222953031452904
depth = 11 Train Score 0.9623451340902863 test Score
0.9252318758281279
depth = 15 Train Score 0.9634267621927706 test Score
0.9231288356496615
depth = 20 Train Score 0.9631629153051491 test Score
0.9235051024711141
depth = 35 Train Score 0.9634333127085721 test Score
0.9235601652753184
depth = 50 Train Score 0.9634333127085721 test Score
0.9235601652753184
depth = 70 Train Score 0.9634333127085721 test Score
0.9235601652753184
depth = 130 Train Score 0.9634333127085721 test Score
0.9235601652753184
```



```
from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
```

```

from scipy.stats import uniform

param_dist = {"n_estimators":sp_randint(105,125),
              "max_depth": sp_randint(10,15),
              "min_samples_split": sp_randint(110,190),
              "min_samples_leaf": sp_randint(25,65)}

clf = RandomForestClassifier(random_state=25,n_jobs=-1)

rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
n_iter=5,cv=10,scoring='f1',random_state=25,return_train_score = True)

rf_random.fit(df_final_train,y_train)
print('mean test scores',rf_random.cv_results_['mean_test_score'])
print('mean train scores',rf_random.cv_results_['mean_train_score'])

mean test scores [0.96225042 0.96215492 0.9605708  0.96194014
0.96330005]
mean train scores [0.96294922 0.96266735 0.96115674 0.96263457
0.96430539]

print(rf_random.cv_results_.keys())

dict_keys(['mean_fit_time', 'std_fit_time', 'mean_score_time',
'std_score_time', 'param_max_depth', 'param_min_samples_leaf',
'param_min_samples_split', 'param_n_estimators', 'params',
'split0_test_score', 'split1_test_score', 'split2_test_score',
'split3_test_score', 'split4_test_score', 'split5_test_score',
'split6_test_score', 'split7_test_score', 'split8_test_score',
'split9_test_score', 'mean_test_score', 'std_test_score',
'rank_test_score', 'split0_train_score', 'split1_train_score',
'split2_train_score', 'split3_train_score', 'split4_train_score',
'split5_train_score', 'split6_train_score', 'split7_train_score',
'split8_train_score', 'split9_train_score', 'mean_train_score',
'std_train_score'])

print(rf_random.best_estimator_)

RandomForestClassifier(max_depth=14, min_samples_leaf=28,
min_samples_split=111,
                    n_estimators=121, n_jobs=-1, random_state=25)

clf = RandomForestClassifier(bootstrap=True, class_weight=None,
criterion='gini',
                    max_depth=14, max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=28, min_samples_split=111,
                    min_weight_fraction_leaf=0.0, n_estimators=121, n_jobs=-1,
                    oob_score=False, random_state=25, verbose=0,
warm_start=False)

```

```

clf.fit(df_final_train,y_train)
y_train_pred = clf.predict(df_final_train)
y_test_pred = clf.predict(df_final_test)

from sklearn.metrics import f1_score
print('Train f1 score',f1_score(y_train,y_train_pred))
print('Test f1 score',f1_score(y_test,y_test_pred))

Train f1 score 0.9652533106548414
Test f1 score 0.9241678239279553

from sklearn.metrics import confusion_matrix
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = (((C.T)/(C.sum(axis=1))).T)

    B =(C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f",
xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f",
xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")


    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f",
xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")

    plt.show()


print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)

```


Train confusion_matrix




Test confusion_matrix



```
from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' %
auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
plt.legend()
plt.show()
```



```
features = df_final_train.columns
importances = clf.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r',
align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```

Assignments:

1. Add another feature called Preferential Attachment with followers and followees data of vertex. you can check about Preferential Attachment in below link
<http://be.amazd.com/link-prediction/>
2. Add feature called svd_dot. you can calculate svd_dot as Dot product between source node svd and destination node svd features. you can read about this in below pdf
https://storage.googleapis.com/kaggle-forum-message-attachments/2594/supervised_link_prediction.pdf
3. Tune hyperparameters for XG boost with all these features and check the error metric.

Adding new features [Preferential Attachment ,SVD_dot]

```
df_final_train.columns
```

```
Index(['jaccard_followers', 'jaccard_followees', 'cosine_followers',
      'cosine_followees', 'num_followers_s', 'num_followees_s',
      'num_followees_d', 'inter_followers', 'inter_followees',
      'adar_index',
      'follows_back', 'same_comp', 'shortest_path', 'weight_in',
      'weight_out',
      'weight_f1', 'weight_f2', 'weight_f3', 'weight_f4',
      'page_rank_s',
      'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hubs_d',
      'authorities_s',
      'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3',
      'svd_u_s_4',
      'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2',
      'svd_u_d_3',
      'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1',
      'svd_v_s_2',
      'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6',
      'svd_v_d_1',
      'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5',
      'svd_v_d_6'],
      dtype='object')
```

Observation:

- source_node,destination_node not found hence we have to get that back in the dataframe to compute new features. hence, reverting data to storage_sample_stage4.

#reading

```
from pandas import read_hdf
df_final_train = read_hdf('storage_sample_stage4.h5',
                          'train_df',mode='r')
df_final_test = read_hdf('storage_sample_stage4.h5',
                         'test_df',mode='r')
```

```
df_final_train.columns
```

```
Index(['source_node', 'destination_node', 'indicator_link',
      'jaccard_followers', 'jaccard_followees', 'cosine_followers',
      'cosine_followees', 'num_followers_s', 'num_followees_s',
      'num_followees_d', 'inter_followers', 'inter_followees',
      'adar_index',
      'follows_back', 'same_comp', 'shortest_path', 'weight_in',
      'weight_out',
      'weight_f1', 'weight_f2', 'weight_f3', 'weight_f4',
      'page_rank_s',
      'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hubs_d',
      'authorities_s',
```

```

        'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3',
'svd_u_s_4',
        'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2',
'svd_u_d_3',
        'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1',
'svd_v_s_2',
        'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6',
'svd_v_d_1',
        'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5',
'svd_v_d_6'],
        dtype='object')

```

```

if os.path.isfile('data/after_eda/train_pos_after_eda.csv'):

```

```

train_graph=nx.read_edgelist('data/after_eda/train_pos_after_eda.csv',
delimiter=',',create_using=nx.DiGraph(),nodetype=int)
print(nx.info(train_graph))

```

```

else:
    print("please run the FB_EDA.ipynb or download the files from
drive")

```

DiGraph with 1780722 nodes and 7550015 edges

```

def followee_preferential_attachment(u1,u2):
    try:
        u_1 = len(set(train_graph.successors(u1)))
        u_2 = len(set(train_graph.successors(u2)))
        return(u_1*u_2)
    except:
        return(0)

```

```

def follower_preferential_attachment(u1,u2):
    try:
        u_1 = len(set(train_graph.predecessors(u1)))
        u_2 = len(set(train_graph.predecessors(u2)))
        return(u_1*u_2)
    except:
        return(0)

```

```

df_final_train.head()

```

	source_node	destination_node	indicator_link	jaccard_followers	\
0	273084	1505602	1	0	
1	832016	1543415	1	0	
2	1325247	760242	1	0	
3	1368400	1006992	1	0	
4	140165	1708748	1	0	

	jaccard_followees	cosine_followers	cosine_followees
num_followers_s \			
0	0.000000	0.000000	0.000000

```

6
1          0.187135          0.028382          0.343828
94
2          0.369565          0.156957          0.566038
28
3          0.000000          0.000000          0.000000
11
4          0.000000          0.000000          0.000000
1

```

```

      num_followees_s  num_followees_d  ...      svd_v_s_3
svd_v_s_4 \
0          15          8  ...  1.983691e-06  1.545075e-13
1          61         142  ... -6.236048e-11  1.345726e-02
2          41          22  ... -2.380564e-19 -7.021227e-19
3           5           7  ...  6.058498e-11  1.514614e-11
4          11           3  ...  1.197283e-07  1.999809e-14

```

```

      svd_v_s_5      svd_v_s_6      svd_v_d_1      svd_v_d_2
svd_v_d_3 \
0  8.108434e-13  1.719702e-14 -1.355368e-12  4.675307e-13  1.128591e-
06
1  3.703479e-12  2.251737e-10  1.245101e-12 -1.636948e-10 -3.112650e-
10
2  1.940403e-19 -3.365389e-19 -1.238370e-18  1.438175e-19 -1.852863e-
19
3  1.513483e-12  4.498061e-13 -9.818087e-10  3.454672e-11  5.213635e-
08
4  3.360247e-13  1.407670e-14  0.000000e+00  0.000000e+00
0.000000e+00

```

```

      svd_v_d_4      svd_v_d_5      svd_v_d_6
0  6.616550e-14  9.771077e-13  4.159752e-14
1  6.738902e-02  2.607801e-11  2.372904e-09
2 -5.901864e-19  1.629341e-19 -2.572452e-19
3  9.595823e-13  3.047045e-10  1.246592e-13
4  0.000000e+00  0.000000e+00  0.000000e+00

```

[5 rows x 54 columns]

```

if not os.path.isfile('data/fea_sample/storage_sample_stage5.h5'):
    df_final_train['followee_preferential_attachment'] =

```

```

df_final_train.apply(lambda row:
followee_preferential_attachment(row['source_node'],row['destination_n
ode']),axis=1)
    df_final_test['followee_preferential_attachment'] =
df_final_test.apply(lambda row:
followee_preferential_attachment(row['source_node'],row['destination_n
ode']),axis=1)

    df_final_train['follower_preferential_attachment'] =
df_final_train.apply(lambda row:
follower_preferential_attachment(row['source_node'],row['destination_n
ode']),axis=1)
    df_final_test['follower_preferential_attachment'] =
df_final_test.apply(lambda row:
follower_preferential_attachment(row['source_node'],row['destination_n
ode']),axis=1)

    hdf = HDFStore('data/fea_sample/storage_sample_stage5.h5')
    hdf.put('train_df',df_final_train, format='table',
data_columns=True)
    hdf.put('test_df',df_final_test, format='table',
data_columns=True)
    hdf.close()
else:
    df_final_train =
read_hdf('data/fea_sample/storage_sample_stage5.h5',
'train_df',mode='r')
    df_final_test =
read_hdf('data/fea_sample/storage_sample_stage5.h5',
'test_df',mode='r')

df_final_train.columns
Index(['source_node', 'destination_node', 'indicator_link',
      'jaccard_followers', 'jaccard_followees', 'cosine_followers',
      'cosine_followees', 'num_followers_s', 'num_followees_s',
      'num_followees_d', 'inter_followers', 'inter_followees',
      'adar_index',
      'follows_back', 'same_comp', 'shortest_path', 'weight_in',
      'weight_out',
      'weight_f1', 'weight_f2', 'weight_f3', 'weight_f4',
      'page_rank_s',
      'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hubs_d',
      'authorities_s',
      'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3',
      'svd_u_s_4',
      'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_2',
      'svd_u_d_3',
      'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_1',

```

```

'svd_v_s_2',
    'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6',
'svd_v_d_1',
    'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_5',
'svd_v_d_6',
    'followee_preferential_attachment',
'follower_preferential_attachment'],
    dtype='object')

```

```
df_final_train.followee_preferential_attachment
```

```

0          120
1         8662
2          902
3           35
4           33

```

```

...
99997       10
99998        4
99999        5
100000       0
100001       0

```

```
Name: followee_preferential_attachment, Length: 100002, dtype: int64
```

Adding SVD Dot feature

#for train datasets

#storing u and v matrices of source and destination nodes in saperate arrays

```

svd_u_s_1,svd_u_s_2,svd_u_s_3,svd_u_s_4,svd_u_s_5,svd_u_s_6 =
df_final_train['svd_u_s_1'],df_final_train['svd_u_s_2'],df_final_train
['svd_u_s_3'],df_final_train['svd_u_s_4'],df_final_train['svd_u_s_5'],
df_final_train['svd_u_s_6']
svd_v_s_1,svd_v_s_2,svd_v_s_3,svd_v_s_4,svd_v_s_5,svd_v_s_6 =
df_final_train['svd_v_s_1'],df_final_train['svd_v_s_2'],df_final_train
['svd_v_s_3'],df_final_train['svd_v_s_4'],df_final_train['svd_v_s_5'],
df_final_train['svd_v_s_6']

```

converting dataframe to np.arrays for dot product computation

```

svd_u_s_1,svd_u_s_2,svd_u_s_3,svd_u_s_4,svd_u_s_5,svd_u_s_6 =
np.array( svd_u_s_1), np.array( svd_u_s_2), np.array( svd_u_s_3),
np.array( svd_u_s_4), np.array( svd_u_s_5), np.array( svd_u_s_6)
svd_v_s_1,svd_v_s_2,svd_v_s_3,svd_v_s_4,svd_v_s_5,svd_v_s_6 =
np.array( svd_v_s_1), np.array( svd_v_s_2), np.array( svd_v_s_3),
np.array( svd_v_s_4), np.array( svd_v_s_5), np.array( svd_v_s_6)

```

```

svd_u_d_1,svd_u_d_2,svd_u_d_3,svd_u_d_4,svd_u_d_5,svd_u_d_6 =
df_final_train['svd_u_d_1'],df_final_train['svd_u_d_2'],df_final_train
['svd_u_d_3'],df_final_train['svd_u_d_4'],df_final_train['svd_u_d_5'],

```

```
df_final_train['svd_u_d_6']
svd_v_d_1,svd_v_d_2,svd_v_d_3,svd_v_d_4,svd_v_d_5,svd_v_d_6 =
df_final_train['svd_v_d_1'],df_final_train['svd_v_d_2'],df_final_train
['svd_v_d_3'],df_final_train['svd_v_d_4'],df_final_train['svd_v_d_5'],
df_final_train['svd_v_d_6']
```

```
svd_u_d_1,svd_u_d_2,svd_u_d_3,svd_u_d_4,svd_u_d_5,svd_u_d_6 =
np.array( svd_u_d_1), np.array( svd_u_d_2), np.array( svd_u_d_3),
np.array( svd_u_d_4), np.array( svd_u_d_5), np.array( svd_u_d_6)
svd_v_d_1,svd_v_d_2,svd_v_d_3,svd_v_d_4,svd_v_d_5,svd_v_d_6 =
np.array( svd_v_d_1), np.array( svd_v_d_2), np.array( svd_v_d_3),
np.array( svd_v_d_4), np.array( svd_v_d_5), np.array( svd_v_d_6)
```

```
svd_u_dot_train=[]
svd_v_dot_train=[]
```

```
for i in range(len(svd_u_s_1)):
```

```
    source_matrix_u=[]
    source_matrix_v=[]
```

```
    destination_matrix_u=[]
    destination_matrix_v=[]
```

```
    source_matrix_u.append(svd_u_s_1[i])
    source_matrix_u.append(svd_u_s_2[i])
    source_matrix_u.append(svd_u_s_3[i])
    source_matrix_u.append(svd_u_s_4[i])
    source_matrix_u.append(svd_u_s_5[i])
    source_matrix_u.append(svd_u_s_6[i])
    source_matrix_v.append(svd_v_s_1[i])
    source_matrix_v.append(svd_v_s_2[i])
    source_matrix_v.append(svd_v_s_3[i])
    source_matrix_v.append(svd_v_s_4[i])
    source_matrix_v.append(svd_v_s_5[i])
    source_matrix_v.append(svd_v_s_6[i])
    destination_matrix_u.append(svd_u_d_1[i])
    destination_matrix_u.append(svd_u_d_2[i])
    destination_matrix_u.append(svd_u_d_3[i])
    destination_matrix_u.append(svd_u_d_4[i])
    destination_matrix_u.append(svd_u_d_5[i])
    destination_matrix_u.append(svd_u_d_6[i])
    destination_matrix_v.append(svd_v_d_1[i])
    destination_matrix_v.append(svd_v_d_2[i])
    destination_matrix_v.append(svd_v_d_3[i])
    destination_matrix_v.append(svd_v_d_4[i])
    destination_matrix_v.append(svd_v_d_5[i])
    destination_matrix_v.append(svd_v_d_6[i])
    #print(source_matrix)
```

```
svd_u_dot_train.append(np.dot(source_matrix_u,destination_matrix_u))
```

```
svd_v_dot_train.append(np.dot(source_matrix_v,destination_matrix_v))
```

```
#for test datasets
```

```
svd_u_s_1,svd_u_s_2,svd_u_s_3,svd_u_s_4,svd_u_s_5,svd_u_s_6 =  
df_final_test['svd_u_s_1'],df_final_test['svd_u_s_2'],df_final_test['s  
vd_u_s_3'],df_final_test['svd_u_s_4'],df_final_test['svd_u_s_5'],df_fi  
nal_test['svd_u_s_6']
```

```
svd_v_s_1,svd_v_s_2,svd_v_s_3,svd_v_s_4,svd_v_s_5,svd_v_s_6 =  
df_final_test['svd_v_s_1'],df_final_test['svd_v_s_2'],df_final_test['s  
vd_v_s_3'],df_final_test['svd_v_s_4'],df_final_test['svd_v_s_5'],df_fi  
nal_test['svd_v_s_6']
```

```
svd_u_s_1,svd_u_s_2,svd_u_s_3,svd_u_s_4,svd_u_s_5,svd_u_s_6 =  
np.array( svd_u_s_1), np.array( svd_u_s_2), np.array( svd_u_s_3),  
np.array( svd_u_s_4), np.array( svd_u_s_5), np.array( svd_u_s_6)  
svd_v_s_1,svd_v_s_2,svd_v_s_3,svd_v_s_4,svd_v_s_5,svd_v_s_6 =  
np.array( svd_v_s_1), np.array( svd_v_s_2), np.array( svd_v_s_3),  
np.array( svd_v_s_4), np.array( svd_v_s_5), np.array( svd_v_s_6)
```

```
svd_u_d_1,svd_u_d_2,svd_u_d_3,svd_u_d_4,svd_u_d_5,svd_u_d_6 =  
df_final_test['svd_u_d_1'],df_final_test['svd_u_d_2'],df_final_test['s  
vd_u_d_3'],df_final_test['svd_u_d_4'],df_final_test['svd_u_d_5'],df_fi  
nal_test['svd_u_d_6']
```

```
svd_v_d_1,svd_v_d_2,svd_v_d_3,svd_v_d_4,svd_v_d_5,svd_v_d_6 =  
df_final_test['svd_v_d_1'],df_final_test['svd_v_d_2'],df_final_test['s  
vd_v_d_3'],df_final_test['svd_v_d_4'],df_final_test['svd_v_d_5'],df_fi  
nal_test['svd_v_d_6']
```

```
svd_u_d_1,svd_u_d_2,svd_u_d_3,svd_u_d_4,svd_u_d_5,svd_u_d_6 =  
np.array( svd_u_d_1), np.array( svd_u_d_2), np.array( svd_u_d_3),  
np.array( svd_u_d_4), np.array( svd_u_d_5), np.array( svd_u_d_6)  
svd_v_d_1,svd_v_d_2,svd_v_d_3,svd_v_d_4,svd_v_d_5,svd_v_d_6 =  
np.array( svd_v_d_1), np.array( svd_v_d_2), np.array( svd_v_d_3),  
np.array( svd_v_d_4), np.array( svd_v_d_5), np.array( svd_v_d_6)
```

```
svd_u_dot_test=[]  
svd_v_dot_test=[]
```

```
for i in range(len(svd_u_s_1)):  
    source_matrix_u=[]  
    source_matrix_v=[]  
  
    destination_matrix_u=[]  
    destination_matrix_v=[]  
  
    source_matrix_u.append(svd_u_s_1[i])  
    source_matrix_u.append(svd_u_s_2[i])
```



```

source_matrix_u.append(svd_u_s_3[i])
source_matrix_u.append(svd_u_s_4[i])
source_matrix_u.append(svd_u_s_5[i])
source_matrix_u.append(svd_u_s_6[i])
source_matrix_v.append(svd_v_s_1[i])
source_matrix_v.append(svd_v_s_2[i])
source_matrix_v.append(svd_v_s_3[i])
source_matrix_v.append(svd_v_s_4[i])
source_matrix_v.append(svd_v_s_5[i])
source_matrix_v.append(svd_v_s_6[i])
destination_matrix_u.append(svd_u_d_1[i])
destination_matrix_u.append(svd_u_d_2[i])
destination_matrix_u.append(svd_u_d_3[i])
destination_matrix_u.append(svd_u_d_4[i])
destination_matrix_u.append(svd_u_d_5[i])
destination_matrix_u.append(svd_u_d_6[i])
destination_matrix_v.append(svd_v_d_1[i])
destination_matrix_v.append(svd_v_d_2[i])
destination_matrix_v.append(svd_v_d_3[i])
destination_matrix_v.append(svd_v_d_4[i])
destination_matrix_v.append(svd_v_d_5[i])
destination_matrix_v.append(svd_v_d_6[i])
#print(source_matrix)

```

```

svd_u_dot_test.append(np.dot(source_matrix_u,destination_matrix_u))

```

```

svd_v_dot_test.append(np.dot(source_matrix_v,destination_matrix_v))

```

```

if not os.path.isfile('data/fea_sample/storage_sample_stage6.h5'):

```

```

    df_final_train['svd_dot_u']=svd_u_dot_train
    df_final_train['svd_dot_v']=svd_v_dot_train
    df_final_test['svd_dot_u']=svd_u_dot_test
    df_final_test['svd_dot_v']=svd_v_dot_test

```

```

    hdf = HDFStore('data/fea_sample/storage_sample_stage6.h5')
    hdf.put('train_df',df_final_train, format='table',
data_columns=True)
    hdf.put('test_df',df_final_test, format='table',
data_columns=True)
    hdf.close()

```

```

else:
    df_final_train =
read_hdf('data/fea_sample/storage_sample_stage6.h5',
'train_df',mode='r')
    df_final_test =
read_hdf('data/fea_sample/storage_sample_stage6.h5',

```

```
'test_df',mode='r')
```

Final features

```
df_final_train.head(4)
```

	source_node	destination_node	indicator_link	jaccard_followers	\
0	273084	1505602	1	0	
1	832016	1543415	1	0	
2	1325247	760242	1	0	
3	1368400	1006992	1	0	

	jaccard_followees	cosine_followers	cosine_followees
num_followers_s \			
0	0.000000	0.000000	0.000000
6			
1	0.187135	0.028382	0.343828
94			
2	0.369565	0.156957	0.566038
28			
3	0.000000	0.000000	0.000000
11			

	num_followees_s	num_followees_d	...	svd_v_d_1
svd_v_d_2 \				
0	15	8	...	-1.355368e-12 4.675307e-13
1	61	142	...	1.245101e-12 -1.636948e-10
2	41	22	...	-1.238370e-18 1.438175e-19
3	5	7	...	-9.818087e-10 3.454672e-11

	svd_v_d_3	svd_v_d_4	svd_v_d_5	svd_v_d_6	\
0	1.128591e-06	6.616550e-14	9.771077e-13	4.159752e-14	
1	-3.112650e-10	6.738902e-02	2.607801e-11	2.372904e-09	
2	-1.852863e-19	-5.901864e-19	1.629341e-19	-2.572452e-19	
3	5.213635e-08	9.595823e-13	3.047045e-10	1.246592e-13	

	followee_preferential_attachment	follower_preferential_attachment
\		
0	120	66
1	8662	1598
2	902	980

```

      svd_dot_u      svd_dot_v
0  1.114958e-11  1.114958e-11
1  3.192812e-03  3.192812e-03
2  1.787503e-35  1.787503e-35
3  4.710376e-20  4.710376e-20

```

[4 rows x 58 columns]

```
df_final_test.head(4)
```

```

      source_node  destination_node  indicator_link  jaccard_followers  \
0          848424          784690             1             0
1          483294          125532             1             0
2          626190          1729265            1             0
3          947219           425228             1             0

```

```

      jaccard_followees  cosine_followers  cosine_followees
num_followers_s  \
0              0.0             0.029161             0.0
14
1              0.0             0.000000             0.0
17
2              0.0             0.000000             0.0
10
3              0.0             0.000000             0.0
37

```

```

      num_followees_s  num_followees_d  ...      svd_v_d_1
svd_v_d_2  \
0              6              9  ... -9.994076e-10  5.791910e-10
1              1              19  ... -9.360516e-12  3.206809e-10
2             16              9  ... -4.253075e-13  4.789463e-13
3             10             34  ... -2.162590e-11  6.939194e-12

```

```

      svd_v_d_3      svd_v_d_4      svd_v_d_5      svd_v_d_6  \
0  3.512364e-07  2.486658e-09  2.771146e-09  1.727694e-12
1  4.668696e-08  6.665777e-12  1.495979e-10  9.836670e-14
2  3.479824e-07  1.630549e-13  3.954708e-13  3.875785e-14
3  1.879861e-05  4.384816e-12  1.239414e-11  6.483485e-13

```

```

      followee_preferential_attachment  follower_preferential_attachment
\

```

0	54	84
1	19	34
2	144	150
3	340	407

	svd_dot_u	svd_dot_v
0	8.425267e-20	2.074808e-17
1	1.352160e-17	1.188376e-17
2	3.671980e-13	3.904885e-12
3	1.634044e-10	9.819784e-11

[4 rows x 58 columns]

```
y_train = df_final_train.indicator_link
y_test = df_final_test.indicator_link
```

```
df_final_train.drop(['source_node',
'destination_node', 'indicator_link'],axis=1,inplace=True)
df_final_test.drop(['source_node',
'destination_node', 'indicator_link'],axis=1,inplace=True)
```

```
df_final_train.head(4)
```

	jaccard_followers	jaccard_followees	cosine_followers
cosine_followees \			
0	0	0.000000	0.000000
0.000000			
1	0	0.187135	0.028382
0.343828			
2	0	0.369565	0.156957
0.566038			
3	0	0.000000	0.000000
0.000000			

	num_followers_s	num_followees_s	num_followees_d	inter_followers
\				
0	6	15	8	0
1	94	61	142	11
2	28	41	22	26
3	11	5	7	0

inter_followees	adar_index	...	svd_v_d_1	svd_v_d_2
-----------------	------------	-----	-----------	-----------

```

svd_v_d_3 \
0      0      0.000000 ... -1.355368e-12  4.675307e-13
1.128591e-06
1      32     16.362912 ...  1.245101e-12 -1.636948e-10 -
3.112650e-10
2      17     10.991826 ... -1.238370e-18  1.438175e-19 -
1.852863e-19
3      0      0.000000 ... -9.818087e-10  3.454672e-11
5.213635e-08

```

```

      svd_v_d_4      svd_v_d_5      svd_v_d_6
followee_preferential_attachment \
0 6.616550e-14 9.771077e-13 4.159752e-14
120
1 6.738902e-02 2.607801e-11 2.372904e-09
8662
2 -5.901864e-19 1.629341e-19 -2.572452e-19
902
3 9.595823e-13 3.047045e-10 1.246592e-13
35

```

```

      follower_preferential_attachment      svd_dot_u      svd_dot_v
0      66 1.114958e-11 1.114958e-11
1      1598 3.192812e-03 3.192812e-03
2      980 1.787503e-35 1.787503e-35
3      22 4.710376e-20 4.710376e-20

```

[4 rows x 55 columns]

Applying XGBoost

```

param_dist = {"n_estimators": sp_randint(105,125),
              "max_depth": sp_randint(2,10),
              "min_child_weight": [2,4,6,8],
              "learning_rate": [0.2,0.4,0.6,0.8]
              }

clf = xgb.XGBClassifier()

rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
n_iter=5,cv=4,scoring='f1',random_state=25,return_train_score = True)

rf_random.fit(df_final_train,y_train)
print('mean test scores',rf_random.cv_results_['mean_test_score'])
print('mean train scores',rf_random.cv_results_['mean_train_score'])

mean test scores [0.98160845 0.98179245 0.98171623 0.98212202
0.98209461]

```

```
mean train scores [0.9965505  0.99976355 1.          0.99886723
0.99990676]
```

```
rf_random.best_estimator_
```

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
              colsample_bylevel=1, colsample_bynode=1,
              colsample_bytree=1,
              early_stopping_rounds=None, enable_categorical=False,
              eval_metric=None, gamma=0, gpu_id=-1,
              grow_policy='depthwise',
              importance_type=None, interaction_constraints='',
              learning_rate=0.4, max_bin=256, max_cat_to_onehot=4,
              max_delta_step=0, max_depth=7, max_leaves=0,
              min_child_weight=8,
              missing=nan, monotone_constraints='()',
              n_estimators=110,
              n_jobs=0, num_parallel_tree=1, predictor='auto',
              random_state=0,
              reg_alpha=0, reg_lambda=1, ...)
```

```
#Best Parameters
```

```
print("Best parameters")
```

```
print("n_estimators :",rf_random.best_estimator_.n_estimators)
print("max_depth: ",rf_random.best_estimator_.max_depth)
print("min_child_weight: ",rf_random.best_estimator_.min_child_weight)
print("learning_rate",rf_random.best_estimator_.learning_rate)
```

```
Best parameters
```

```
n_estimators : 110
max_depth: 7
min_child_weight: 8
learning_rate 0.4
```

```
# using best parameters
```

```
clf=xgb.XGBClassifier(base_score=0.5, booster='gbtree',
                      callbacks=None,
                      colsample_bylevel=1, colsample_bynode=1,
                      colsample_bytree=1,
                      early_stopping_rounds=None, enable_categorical=False,
                      eval_metric=None, gamma=0, gpu_id=-1,
                      grow_policy='depthwise',
                      importance_type=None, interaction_constraints='',
                      learning_rate=0.4, max_bin=256, max_cat_to_onehot=4,
                      max_delta_step=0, max_depth=7, max_leaves=0,
                      min_child_weight=8,
                      monotone_constraints='()', n_estimators=110,
                      n_jobs=0, num_parallel_tree=1, predictor='auto',
```

```

random_state=0,
                reg_alpha=0, reg_lambda=1)

clf.fit(df_final_train,y_train)
y_train_pred = clf.predict(df_final_train)
y_test_pred = clf.predict(df_final_test)

print('Train f1 score',f1_score(y_train,y_train_pred))
print('Test f1 score',f1_score(y_test,y_test_pred))

Train f1 score 0.9981308037543857
Test f1 score 0.9175332734197801

def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = (((C.T)/(C.sum(axis=1))).T)

    B =(C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f",
xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f",
xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f",
xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")


    plt.show()

print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)


```

```
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)
```


Train confusion_matrix




Test confusion_matrix



```
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' %
auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
plt.legend()
plt.show()
```

```
features = df_final_train.columns
importances = clf.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r',
align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



Obervation:

- Preferential attachment and svd_dot features found to be not important as per XGBoost model
- XGBoost performs similar to RandomForest in results but from performace point of view XGBoost took longer duration in hyper-paramter tuning